

04b_data_pandas

November 23, 2022

1 Doing More with Data: pandas

1.1 Introduction to Python

Data Sciences Institute, University of Toronto

Instructor: Kaylie Lau | TA: Salaar Liaqat

November - December 2022

2 Contents:

1. Setup
2. Intro to pandas
3. Getting data
4. Profiling and initial data exploration: changing data types, descriptive statistics
5. Wrangling and plotting: concatenating, merging, adding and removing columns, filtering and selecting, null values, grouping and aggregating, plotting
6. Writing to file
7. More wrangling: reshaping, applying functions

2.1 Data

This module uses four datasets: [bike thefts](#), [TTC subway delays](#) and [subway delay reason codes](#), and [neighbourhood profiles](#). All four are available in the course repo, and originally come from Toronto Open Data.

The specific file names are: - bicycle-thefts - 4326.csv - ttc-subway-delay-data-2021.xlsx - ttc-subway-delay-codes.xlsx - neighbourhood-profiles-2016-140-model.csv

3 pandas

3.1 What is pandas?

pandas is a package for data analysis and manipulation. (The name is a reference to panel data, not the animal.) It gives us data frames, which represent data in a table of columns and rows, and

functions to manipulate and plot them. `pandas` also provides a slew of functions for reading and writing data to a variety of sources, including files, SQL databases, and compressed binary formats.

```
[ ]: import numpy as np
      # pd is the conventional alias for pandas
      import pandas as pd

      # display all columns
      pd.set_option("display.max_columns", None)
```

3.2 DataFrames

Columns are labeled with their names. Rows also have a label, or *index*. If row labels are not specified, `pandas` uses numbers as the default. Each column is a *Series*, or one-dimensional array, where values share a data type. Unlike `numpy` arrays, DataFrames can have columns of different data types. However, like arrays and lists, **DataFrames are mutable** -- this means that if more than one variable refers to the same DataFrame, updating one updates them all!

3.3 Getting data

We can create a DataFrame manually with `DataFrame()` constructor. If a dictionary is passed to `DataFrame()`, the keys become column names, and the values become the rows. Calling just `DataFrame()` creates an empty DataFrame to which data can be added later.

```
[ ]: trees = pd.DataFrame({
      'name': ['sugar maple', 'black oak', 'white ash', 'douglas fir'],
      'avg_lifespan': [300, 100, 260, 450],
      'quantity': [53, 207, 178, 93]
    })
trees
```

```
[ ]:      name  avg_lifespan  quantity
0  sugar maple           300         53
1   black oak            100        207
2   white ash            260        178
3  douglas fir            450         93
```

We can create an individual column with `Series()`. The `name` argument corresponds to a column name.

```
[ ]: tree_types = pd.Series(['deciduous', 'deciduous', 'deciduous', 'evergreen'],
                             name='foliage')
tree_types
```

```
[ ]: 0    deciduous
     1    deciduous
     2    deciduous
```

```
3    evergreen
Name: foliage, dtype: object
```

3.3.1 Data from csv

Of course, we're more likely to load data into a DataFrame than to create DataFrames manually. `pandas` has read functions for different file formats. To read data from a csv or other delimited file, we use `pd.read_csv()`, then pass in the local file path or the URL of the csv to read. `pandas` will infer the data type of each column based on the values in the first chunk of the file loaded.

```
[ ]: thefts = pd.read_csv('/content/data/bicycle-thefts - 4326.csv')
```

3.4 Profiling and initial data cleaning

We got our data, but now we need to understand what's in it. We can start to understand the DataFrame by checking out its `dtypes` and `shape` attributes, which give column data types and row by column dimensions, respectively. Note that `object` is `pandas`' way of saying values are represented as string data.

```
[ ]: thefts.shape
```

```
[ ]: (25569, 33)
```

```
[ ]: thefts.dtypes
```

```
[ ]: _id                int64
    OBJECTID            int64
    event_unique_id      object
    Primary_Offence       object
    Occurrence_Date       object
    Occurrence_Year       int64
    Occurrence_Month      object
    Occurrence_DayOfWeek  object
    Occurrence_DayOfMonth int64
    Occurrence_DayOfYear  int64
    Occurrence_Hour       int64
    Report_Date           object
    Report_Year           int64
    Report_Month          object
    Report_DayOfWeek      object
    Report_DayOfMonth     int64
    Report_DayOfYear      int64
    Report_Hour           int64
    Division             object
    City                 object
```

```

Hood_ID                object
NeighbourhoodName      object
Location_Type          object
Premises_Type          object
Bike_Make              object
Bike_Model             object
Bike_Type              object
Bike_Speed             int64
Bike_Colour            object
Cost_of_Bike           float64
Status                 object
ObjectId2              int64
geometry               object
dtype: object

```

3.4.1 head()s and tail()s

To check out the first few rows, we can call the DataFrame `head()` method. Similarly, we can see the last few rows with the `tail()` method. Five rows are shown by default, but we can change that by passing an integer as an argument.

```
[ ]: thefts.head()
```

```

[ ]:
   _id  OBJECTID event_unique_id      Primary_Offence  Occurrence_Date \
0    1    17744  GO-20179016397          THEFT UNDER  2017-10-03T00:00:00
1    2    17759  GO-20172033056  THEFT UNDER - BICYCLE  2017-11-08T00:00:00
2    3    17906  GO-20189030822  THEFT UNDER - BICYCLE  2018-09-14T00:00:00
3    4    17962   GO-2015804467          THEFT UNDER  2015-05-07T00:00:00
4    5    17963  GO-20159002781          THEFT UNDER  2015-05-16T00:00:00

   Occurrence_Year Occurrence_Month Occurrence_DayOfWeek \
0              2017           October           Tuesday
1              2017           November          Wednesday
2              2018           September            Friday
3              2015              May            Thursday
4              2015              May            Saturday

   Occurrence_DayOfMonth Occurrence_DayOfYear Occurrence_Hour \
0                      3                276             14
1                      8                312              3
2                     14                257              9
3                      7                127             18
4                     16                136             12

   Report_Date Report_Year Report_Month Report_DayOfWeek \
0  2017-10-03T00:00:00      2017     October           Tuesday

```

1	2017-11-08T00:00:00	2017	November	Wednesday
2	2018-09-17T00:00:00	2018	September	Monday
3	2015-05-14T00:00:00	2015	May	Thursday
4	2015-05-16T00:00:00	2015	May	Saturday

	Report_DayOfMonth	Report_DayOfYear	Report_Hour	Division	City	Hood_ID	\
0	3	276	18	D22	Toronto	15	
1	8	312	22	D22	Toronto	15	
2	17	260	16	D22	Toronto	15	
3	14	134	14	D22	Toronto	15	
4	16	136	15	D22	Toronto	15	

	NeighbourhoodName	Location_Type	\
0	Kingsway South (15)	Streets, Roads, Highways (Bicycle Path, Privat...	
1	Kingsway South (15)	Single Home, House (Attach Garage, Cottage, Mo...	
2	Kingsway South (15)	Ttc Subway Station	
3	Kingsway South (15)	Ttc Subway Station	
4	Kingsway South (15)	Ttc Subway Station	

	Premises_Type	Bike_Make	Bike_Model	Bike_Type	Bike_Speed	Bike_Colour	\
0	Outside	GI	ESCAPE 2	OT	7	BLK	
1	House	UNKNOWN	MAKE	NaN	1	BLK	
2	Transit	OT	CROSSTAIL	MT	24	BLK	
3	Transit	GT	NaN	TO	10	BLKDGR	
4	Transit	GI	NaN	MT	6	RED	

	Cost_of_Bike	Status	ObjectId2	\
0	700.0	STOLEN	1	
1	1100.0	RECOVERED	2	
2	904.0	STOLEN	3	
3	400.0	STOLEN	4	
4	600.0	STOLEN	5	

	geometry
0	{'type': 'Point', 'coordinates': (-79.50655965...
1	{'type': 'Point', 'coordinates': (-79.50484874...
2	{'type': 'Point', 'coordinates': (-79.51170915...
3	{'type': 'Point', 'coordinates': (-79.51170915...
4	{'type': 'Point', 'coordinates': (-79.51132657...

```
[ ]: # last 3
thefts.tail(3)
```

	_id	OBJECTID	event_unique_id	Primary_Offence	\
25566	25567	11462	G0-20169005434	THEFT UNDER	
25567	25568	11695	G0-20161170896	THEFT UNDER	
25568	25569	11883	G0-20169007653	THEFT UNDER - BICYCLE	

	Occurrence_Date	Occurrence_Year	Occurrence_Month	\
25566	2016-06-04T00:00:00	2016	June	
25567	2016-07-04T00:00:00	2016	July	
25568	2016-07-22T00:00:00	2016	July	

	Occurrence_DayOfWeek	Occurrence_DayOfMonth	Occurrence_DayOfYear	\
25566	Saturday	4	156	
25567	Monday	4	186	
25568	Friday	22	204	

	Occurrence_Hour	Report_Date	Report_Year	Report_Month	\
25566	22	2016-06-07T00:00:00	2016	June	
25567	20	2016-07-04T00:00:00	2016	July	
25568	9	2016-07-23T00:00:00	2016	July	

	Report_DayOfWeek	Report_DayOfMonth	Report_DayOfYear	Report_Hour	\
25566	Tuesday	7	159	16	
25567	Monday	4	186	20	
25568	Saturday	23	205	11	

	Division	City	Hood_ID	NeighbourhoodName	\
25566	D42	Toronto	132	Malvern (132)	
25567	D42	Toronto	132	Malvern (132)	
25568	D42	Toronto	132	Malvern (132)	

	Location_Type	Premises_Type	\
25566	Apartment (Rooming House, Condo)	Apartment	
25567	Other Commercial / Corporate Places (For Profi...	Commercial	
25568	Parking Lots (Apt., Commercial Or Non-Commercial)	Outside	

	Bike_Make	Bike_Model	Bike_Type	Bike_Speed	Bike_Colour	\
25566	SC	ANTRIM	MT	24	WHI	
25567	UNKNOWN MAKE	NaN	SC	1	NaN	
25568	SU	ASCENT MOUNTAIN	MT	21	ONG	

	Cost_of_Bike	Status	ObjectId2	\
25566	700.0	STOLEN	25567	
25567	3000.0	STOLEN	25568	
25568	200.0	STOLEN	25569	

	geometry
25566	{'type': 'Point', 'coordinates': (-79.2360175,...
25567	{'type': 'Point', 'coordinates': (-79.20060719...
25568	{'type': 'Point', 'coordinates': (-79.23734742...

3.4.2 Renaming columns

Most, but not all, of the bike theft columns follow the same naming convention. For convenience's sake, though, let's convert the column names to all lowercase. We can do this with the DataFrame `rename()` method. `rename()` accepts either a dictionary with current column names as the keys and new names as the values, or the name of a function to transform names. Let's write a function.

```
[ ]: # notice that we do not add () to the function name
thefts = thefts.rename(columns=str.lower)
```

Let's also rename `cost_of_bike` so it follows the pattern of the other bike attribute columns.

```
[ ]: thefts = thefts.rename(columns={'cost_of_bike': 'bike_cost'})

# view column names
print(list(thefts))
```

```
['_id', 'objectid', 'event_unique_id', 'primary_offence', 'occurrence_date',
'occurrence_year', 'occurrence_month', 'occurrence_dayofweek',
'occurrence_dayofmonth', 'occurrence_dayofyear', 'occurrence_hour',
'report_date', 'report_year', 'report_month', 'report_dayofweek',
'report_dayofmonth', 'report_dayofyear', 'report_hour', 'division', 'city',
'hood_id', 'neighbourhoodname', 'location_type', 'premises_type', 'bike_make',
'bike_model', 'bike_type', 'bike_speed', 'bike_colour', 'bike_cost', 'status',
'objectid2', 'geometry']
```

3.4.3 Profiling columns

It can be useful to focus on a subset of columns, particularly to understand value sets. To select a single column in a DataFrame, we can supply the name of the column in square brackets, just like we did when accessing values in a dictionary. `pandas` will return the column as a Series. To get unique values, we can use the `unique()` Series method. If we want to count how many times each value appears, we can use the `value_counts()` method.

```
[ ]: thefts['status']
```

```
[ ]: 0          STOLEN
     1    RECOVERED
     2          STOLEN
     3          STOLEN
     4          STOLEN
     ...
    25564    STOLEN
    25565    STOLEN
    25566    STOLEN
    25567    STOLEN
    25568    STOLEN
```

Name: status, Length: 25569, dtype: object

```
[ ]: thefts['status'].unique()
```

```
[ ]: array(['STOLEN', 'RECOVERED', 'UNKNOWN'], dtype=object)
```

```
[ ]: thefts['status'].value_counts()
```

```
[ ]: STOLEN      24807
      UNKNOWN    454
      RECOVERED  308
      Name: status, dtype: int64
```

We can summarize numeric Series much like we did with `numpy` functions.

```
[ ]: thefts['bike_cost'].median()
```

```
[ ]: 600.0
```

```
[ ]: thefts['bike_cost'].quantile(0.9)
```

```
[ ]: 2000.0
```

3.4.4 `info()`

We can get an overview of the `DataFrame` by profiling it with the `info()` method.

`info()` prints a lot of information about a `DataFrame`, including:

- * the **shape** as the number of rows and columns
- * column names and their **dtype**
- * the number of non-null values in each column
- * how big the `DataFrame` is in terms of memory usage

The bicycle theft data looks quite complete, though some records are missing bike descriptors like `bike_make`, `bike_model`, `bike_colour`, and `bike_cost`.

Most of the column dtypes make sense. We'll want to convert the dates to proper dates. We may also want to convert string columns with limited value sets, like `status`, to categorical data.

```
[ ]: thefts.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25569 entries, 0 to 25568
Data columns (total 33 columns):
#   Column                Non-Null Count  Dtype
---  -
0   _id                    25569 non-null  int64
1   objectid               25569 non-null  int64
2   event_unique_id        25569 non-null  object
3   primary_offence        25569 non-null  object
```



```

4  occurrence_date      25569 non-null object
5  occurrence_year      25569 non-null int64
6  occurrence_month     25569 non-null object
7  occurrence_dayofweek  25569 non-null object
8  occurrence_dayofmonth 25569 non-null int64
9  occurrence_dayofyear  25569 non-null int64
10 occurrence_hour      25569 non-null int64
11 report_date          25569 non-null object
12 report_year          25569 non-null int64
13 report_month         25569 non-null object
14 report_dayofweek     25569 non-null object
15 report_dayofmonth    25569 non-null int64
16 report_dayofyear     25569 non-null int64
17 report_hour          25569 non-null int64
18 division             25569 non-null object
19 city                 25569 non-null object
20 hood_id              25569 non-null object
21 neighbourhoodname    25569 non-null object
22 location_type        25569 non-null object
23 premises_type        25569 non-null object
24 bike_make            25448 non-null object
25 bike_model           15923 non-null object
26 bike_type            25569 non-null object
27 bike_speed           25569 non-null int64
28 bike_colour          23508 non-null object
29 bike_cost            23825 non-null float64
30 status               25569 non-null object
31 objectid2           25569 non-null int64
32 geometry             25569 non-null object
dtypes: float64(1), int64(12), object(20)
memory usage: 6.4+ MB

```

3.4.5 Changing data types

Before exploring the bike theft data further, let's fix the date and categorical columns. To convert a column to datetime, we use the `pd.to_datetime()` function, passing in the column to convert, and reassign the output back to the column we're converting.

pandas knows how to convert the dates in the bike thefts data, but for less common formats, it is necessary to use the `format` keyword argument to specify how dates should be parsed. `format` strings use `strftime` codes. See <https://strftime.org/> for a cheat sheet.

```
[ ]: thefts['occurrence_date'] = pd.to_datetime(thefts['occurrence_date'],
                                              format='%Y-%m-%d')
thefts['occurrence_date']
```

```
[ ]: 0      2017-10-03
      1      2017-11-08
      2      2018-09-14
      3      2015-05-07
      4      2015-05-16
      ...
      25564   2015-04-01
      25565   2016-05-16
      25566   2016-06-04
      25567   2016-07-04
      25568   2016-07-22
      Name: occurrence_date, Length: 25569, dtype: datetime64[ns]
```

```
[ ]: # convert report_date without the format argument
      thefts['report_date'] = pd.to_datetime(thefts['report_date'])
      thefts['report_date']
```

```
[ ]: 0      2017-10-03
      1      2017-11-08
      2      2018-09-17
      3      2015-05-14
      4      2015-05-16
      ...
      25564   2015-04-01
      25565   2016-05-16
      25566   2016-06-07
      25567   2016-07-04
      25568   2016-07-23
      Name: report_date, Length: 25569, dtype: datetime64[ns]
```

All other data type conversions can be done with the `astype()` method. If we were converting to a number, `pd.to_numeric()` provides an easy way to convert without having to pick a specific numeric data type.

```
[ ]: thefts['status'] = thefts['status'].astype('category')
      thefts['status']
```

```
[ ]: 0      STOLEN
      1    RECOVERED
      2      STOLEN
      3      STOLEN
      4      STOLEN
      ...
      25564   STOLEN
      25565   STOLEN
      25566   STOLEN
      25567   STOLEN
```

```

25568      STOLEN
Name: status, Length: 25569, dtype: category
Categories (3, object): ['RECOVERED', 'STOLEN', 'UNKNOWN']

```

We can select and convert multiple columns at once by passing a list of columns in the square brackets., then using `.astype()`.

```

[ ]: thefts[['location_type', 'premises_type']] =
    ↳ thefts[['location_type', 'premises_type']].astype('category')

# check data types
thefts[['location_type', 'premises_type']].dtypes

```

```

[ ]: location_type    category
     premises_type    category
     dtype: object

```

3.4.6 describe()

To get a sense of the values in a DataFrame, we can use the `describe()` method. `describe()` summarizes only numeric columns by default. Passing the `include='all'` argument will produce summary statistics for other columns as well.

```

[ ]: thefts.describe(include='all',
                      datetime_is_numeric=True) # silence warning about upcoming
    ↳ pandas change

```

```

[ ]:
count      25569.000000  25569.000000  25569  25569
unique           NaN           NaN  22771    66
top           NaN           NaN  GO-20201550944  THEFT UNDER
freq           NaN           NaN    14  11904
mean    12785.000000  12909.173218    NaN    NaN
min         1.000000    1.000000    NaN    NaN
25%         6393.000000    6456.000000    NaN    NaN
50%    12785.000000  12918.000000    NaN    NaN
75%    19177.000000  19360.000000    NaN    NaN
max     25569.000000  25806.000000    NaN    NaN
std       7381.278853    7448.318562    NaN    NaN

              occurrence_date  occurrence_year  occurrence_month  \
count              25569      25569.000000      25569
unique              NaN              NaN              12
top              NaN              NaN              July
freq              NaN              NaN              4002
mean    2017-09-04 03:39:28.321013504      2017.124174      NaN

```

min	2009-09-01 00:00:00	2009.000000	NaN
25%	2016-01-06 00:00:00	2016.000000	NaN
50%	2017-09-05 00:00:00	2017.000000	NaN
75%	2019-06-20 00:00:00	2019.000000	NaN
max	2020-12-30 00:00:00	2020.000000	NaN
std	NaN	1.960127	NaN

	occurrence_dayofweek	occurrence_dayofmonth	occurrence_dayofyear	\
count	25569	25569.000000	25569.000000	
unique	7	NaN	NaN	
top	Friday	NaN	NaN	
freq	3924	NaN	NaN	
mean	NaN	15.616684	202.227698	
min	NaN	1.000000	1.000000	
25%	NaN	8.000000	153.000000	
50%	NaN	16.000000	205.000000	
75%	NaN	23.000000	259.000000	
max	NaN	31.000000	366.000000	
std	NaN	8.592886	76.821431	

	occurrence_hour	report_date	report_year	\
count	25569.000000	25569	25569.000000	
unique	NaN	NaN	NaN	
top	NaN	NaN	NaN	
freq	NaN	NaN	NaN	
mean	13.274395	2017-09-12 12:02:37.127771904	2017.143572	
min	0.000000	2014-01-01 00:00:00	2014.000000	
25%	9.000000	2016-01-22 00:00:00	2016.000000	
50%	14.000000	2017-09-12 00:00:00	2017.000000	
75%	19.000000	2019-06-26 00:00:00	2019.000000	
max	23.000000	2020-12-31 00:00:00	2020.000000	
std	6.530181	NaN	1.955024	

	report_month	report_dayofweek	report_dayofmonth	report_dayofyear	\
count	25569	25569	25569.000000	25569.000000	
unique	12	7	NaN	NaN	
top	July	Monday	NaN	NaN	
freq	3988	4318	NaN	NaN	
mean	NaN	NaN	15.924870	203.493723	
min	NaN	NaN	1.000000	1.000000	
25%	NaN	NaN	9.000000	154.000000	
50%	NaN	NaN	16.000000	206.000000	
75%	NaN	NaN	23.000000	260.000000	
max	NaN	NaN	31.000000	366.000000	
std	NaN	NaN	8.549584	77.115977	

report_hour	division	city	hood_id	\
-------------	----------	------	---------	---

count	25569.000000	25569	25569	25569
unique	NaN	18	2	141
top	NaN	D14	Toronto	77
freq	NaN	4580	25560	2576
mean	14.224139	NaN	NaN	NaN
min	0.000000	NaN	NaN	NaN
25%	11.000000	NaN	NaN	NaN
50%	14.000000	NaN	NaN	NaN
75%	18.000000	NaN	NaN	NaN
max	23.000000	NaN	NaN	NaN
std	5.052944	NaN	NaN	NaN

	neighbourhoodname	\
count		25569
unique		141
top	Waterfront Communities-The Island	(77)
freq		2576
mean		NaN
min		NaN
25%		NaN
50%		NaN
75%		NaN
max		NaN
std		NaN

	location_type	premises_type	bike_make	bike_model	\
count		25569	25569	25448	15923
unique		42	7	820	8097
top	Apartment (Rooming House, Condo)		Outside	OT	UNKNOWN
freq		5887	7960	4991	304
mean		NaN	NaN	NaN	NaN
min		NaN	NaN	NaN	NaN
25%		NaN	NaN	NaN	NaN
50%		NaN	NaN	NaN	NaN
75%		NaN	NaN	NaN	NaN
max		NaN	NaN	NaN	NaN
std		NaN	NaN	NaN	NaN

	bike_type	bike_speed	bike_colour	bike_cost	status	\
count	25569	25569.000000	23508	23825.000000	25569	
unique	13	NaN	252	NaN	3	
top	MT	NaN	BLK	NaN	STOLEN	
freq	8245	NaN	7422	NaN	24807	
mean	NaN	14.164144	NaN	949.542371	NaN	
min	NaN	0.000000	NaN	0.000000	NaN	
25%	NaN	6.000000	NaN	350.000000	NaN	
50%	NaN	15.000000	NaN	600.000000	NaN	

75%	NaN	21.000000	NaN	1000.000000	NaN
max	NaN	99.000000	NaN	120000.000000	NaN
std	NaN	10.559215	NaN	1675.880345	NaN

	objectid2	geometry
count	25569.000000	25569
unique	NaN	5816
top	NaN	{'type': 'Point', 'coordinates': (-79.38372586...
freq	NaN	167
mean	12785.000000	NaN
min	1.000000	NaN
25%	6393.000000	NaN
50%	12785.000000	NaN
75%	19177.000000	NaN
max	25569.000000	NaN
std	7381.278853	NaN

3.5 Wrangling and Plotting

3.5.1 Combining datasets: concatenation

Just as `pandas` has `read_csv()` for flat files, there is a `read_excel()` function to load Excel files.

The TTC publishes subway delay data as a multi-sheet Excel workbook, with a month's worth of data per sheet. `read_excel()` loads just the first sheet in an Excel file by default. To load all sheets, pass in the keyword argument `sheet_name=None`. The result is a dictionary, where each key is the sheet name and each value is a `DataFrame` with the contents of the sheet.

```
[ ]: delays = pd.read_excel('/content/data/ttc-subway-delay-data-2021.xlsx',
    ↪sheet_name=None)
```

```
[ ]: type(delays)
```

```
[ ]: dict
```

To combine them all, we create an empty `DataFrame`, then loop through the dictionary items and use `pd.concat()` to append data. `concat()` takes a list of `DataFrames` to combine. Since we did not specify an index, row labels are numbers: the first row of each sheet has an index of 0, and so on. To reset row labels so that they are sequential again, we set `ignore_index=True`.

```
[ ]: # create an empty DataFrame
all_delays = pd.DataFrame()

for sheet_name, values in delays.items():
    # print the number of rows
    print(f'Adding {values.shape[0]} rows from {sheet_name}')
    # add each sheet to all_delays
```

```

all_delays = pd.concat([all_delays, values],
                        axis=0, # concatenate rows
                        ignore_index=True) # reset row labels

all_delays.shape

```

```

Adding 1216 rows from January21
Adding 1245 rows from Feb 21
Adding 1167 rows from March '21
Adding 1170 rows from April '21
Adding 1168 rows from May '21
Adding 1265 rows from June 21
Adding 1244 rows from July 21
Adding 1273 rows from August 21
Adding 1433 rows from Sept 21
Adding 1560 rows from Oct 21
Adding 1771 rows from Nov 21
Adding 1858 rows from December21

```

```
[ ]: (16370, 10)
```

```
[ ]: all_delays.head()
```

```
[ ]:
      Date   Time   Day      Station   Code   Min Delay   Min Gap \
0 2021-01-01 00:33 Friday    BLOOR STATION  MUPAA         0         0
1 2021-01-01 00:39 Friday  SHERBOURNE STATION  EUCO          5         9
2 2021-01-01 01:07 Friday  KENNEDY BD STATION  EUCD          5         9
3 2021-01-01 01:41 Friday    ST CLAIR STATION  MUIS          0         0
4 2021-01-01 02:04 Friday  SHEPPARD WEST STATION  MUIS          0         0

      Bound Line  Vehicle
0         N    YU      6046
1         E    BD      5250
2         E    BD      5249
3        NaN    YU         0
4        NaN    YU         0

```

The TTC delays data includes a reason code for the delay. Code definitions, however, are in a separate Excel file, `ttc-subway-delay-codes.xlsx`. This file has been modified slightly to make it easier to work with. Codes are split between two tabs, so we will load both to a DataFrame, `delay_reasons`.

```

[ ]: dr = pd.read_excel('/content/data/ttc-subway-delay-codes.xlsx', sheet_name=None)

delay_reasons = pd.DataFrame()
for sheet_name, values in dr.items():
    delay_reasons = pd.concat([delay_reasons, values],

```

```
axis=0,
ignore_index=True)
```

```
delay_reasons
```

```
[ ]:
RMENU CODE          CODE DESCRIPTION SUB OR SRT
0      EUAC          Air Conditioning      SUB
1      EUAL          Alternating Current    SUB
2      EUATC         ATC RC&S Equipment     SUB
3      EUBK          Brakes                 SUB
4      EUBO          Body                   SUB
..      ...
195    TRNOA         No Operator Immediately Available SRT
196    TRO          Transportation Department - Other SRT
197    TRSET        Train Controls Improperly Shut Down SRT
198    TRST          Storm Trains           SRT
199    TRTC         Transit Control Related Problems SRT
```

```
[200 rows x 3 columns]
```

We will rename the columns in both `all_delays` and `delay_reasons` so that we replace spaces with underscores as well as convert all letters to lowercase.

```
[ ]: def clean_names(string):
      return string.lower().replace(' ', '_')

delay_reasons = delay_reasons.rename(columns=clean_names)
all_delays = all_delays.rename(columns=clean_names)
print(list(delay_reasons))
print(list(all_delays))
```

```
['rmenu_code', 'code_description', 'sub_or_srt']
['date', 'time', 'day', 'station', 'code', 'min_delay', 'min_gap', 'bound',
'line', 'vehicle']
```

3.6 Combining datasets: merging

Ideally, the delays data would include code descriptions. We can get descriptions into `all_delays` by *merging* in `delay_reasons`. Merging is analogous to joining in SQL databases. To merge two DataFrames, we pass them as arguments to the `pd.merge()`. Then, we specify how to merge the two DataFrames and what column names to merge on.

Let's review the `all_delays` and `delay_reasons` DataFrames. `code` is equivalent to `rmenu_code`. If we pass in `all_delays` as the first DataFrame, then it will be the left frame, and `delay_reasons` the right one. We want to keep all the delay records, even if there isn't a matching code in `delay_reasons`, so we will perform a left join.

```
[ ]: all_delays.head(2)
```



```
[ ]:      date    time    day          station  code  min_delay  min_gap  \
0 2021-01-01  00:33  Friday      BLOOR STATION  MUPAA          0         0
1 2021-01-01  00:39  Friday  SHERBOURNE STATION  EUCO            5         9

      bound line  vehicle
0         N   YU    6046
1         E   BD    5250
```

```
[ ]: delay_reasons.head(2)
```

```
[ ]:  rmenu_code    code_description  sub_or_srt
0         EUAC      Air Conditioning      SUB
1         EUAL  Alternating Current      SUB
```

```
[ ]: delays_w_reasons = pd.merge(all_delays,
                                delay_reasons,
                                how='left',
                                left_on='code',
                                right_on='rmenu_code')
delays_w_reasons.head(3)
```

```
[ ]:      date    time    day          station  code  min_delay  min_gap  \
0 2021-01-01  00:33  Friday      BLOOR STATION  MUPAA          0         0
1 2021-01-01  00:39  Friday  SHERBOURNE STATION  EUCO            5         9
2 2021-01-01  01:07  Friday  KENNEDY BD STATION  EUCD            5         9

      bound line  vehicle  rmenu_code  \
0         N   YU    6046      MUPAA
1         E   BD    5250      EUCO
2         E   BD    5249      EUCD

                                code_description  sub_or_srt
0  Passenger Assistance Alarm Activated - No Trou...      SUB
1                                Couplers              SUB
2      Consequential Delay (2nd Delay Same Fault)      SUB
```

3.7 drop()

The resulting DataFrame has both our join columns, which is redundant. We can drop one with the `drop()` DataFrame method, passing in the column name(s) we want to drop in the `columns` keyword argument.

```
[ ]: delays_w_reasons = delays_w_reasons.drop(columns='rmenu_code')
delays_w_reasons.head(3)
```

```
[ ]:      date    time    day      station    code  min_delay  min_gap  \
0 2021-01-01  00:33  Friday      BLOOR STATION  MUPAA           0         0
1 2021-01-01  00:39  Friday  SHERBOURNE STATION  EUCO            5         9
2 2021-01-01  01:07  Friday  KENNEDY BD STATION  EUCD            5         9

      bound line  vehicle      code_description  \
0      N   YU    6046  Passenger Assistance Alarm Activated - No Trou...
1      E   BD    5250                        Couplers
2      E   BD    5249      Consequential Delay (2nd Delay Same Fault)

      sub_or_srt
0          SUB
1          SUB
2          SUB
```

3.8 Creating new columns

Adding a column to a DataFrame looks like adding a key-value pair to a dictionary. At its simplest, we can assign a single value to repeat down a column.

```
[ ]: delays_w_reasons['year'] = 2021
delays_w_reasons['year'].unique()
```

```
[ ]: array([2021])
```

We can also write an expression and store the resulting values in a new column.

```
[ ]: delays_w_reasons['hour_delay'] = round(delays_w_reasons['min_delay'] / 60, 2)
delays_w_reasons[['min_delay', 'hour_delay']].head()
```

```
[ ]:      min_delay  hour_delay
0         0         0.00
1         5         0.08
2         5         0.08
3         0         0.00
4         0         0.00
```

It is also possible to extract parts of a datetime column with the `dt` accessor.

```
[ ]: delays_w_reasons['month'] = delays_w_reasons['date'].dt.month
delays_w_reasons['month']
```

```
[ ]: 0      1
1      1
2      1
3      1
4      1
```

```

..
16365    12
16366    12
16367    12
16368    12
16369    12
Name: month, Length: 16370, dtype: int64

```

It is possible to create a new integer column, `hour`, that contains the hour in which a delay occurred. Below we highlight two methods.

```

[ ]: # two ways to extract hour
     # convert to time, then access hour
delays_w_reasons['hour'] = pd.to_datetime(delays_w_reasons['time'], format='%H:
     ↳%M').dt.hour

     # split and take first part
delays_w_reasons['hour'] = delays_w_reasons['time'].str.split(':',
     ↳expand=True)[0]
delays_w_reasons['hour'] = delays_w_reasons['hour'].astype(int)

```

3.9 Filtering and selecting data

Let's take another look at the TTC subway delay data. There are only 4 subway lines in Toronto, but `describe()` reported 17 unique values.

```

[ ]: delays_w_reasons['line'].unique()

[ ]: array(['YU', 'BD', 'SHP', 'SRT', 'YU/BD', nan, 'YONGE/UNIVERSITY/BLOOR',
          'YU / BD', 'YUS', '999', 'SHEP', '36 FINCH WEST', 'YUS & BD',
          'YU & BD LINES', '35 JANE', '52', '41 KEELE', 'YUS/BD'],
          dtype=object)

```

Looks like some of the line values should be updated (YU/BD variants) and others should be dropped (e.g., 36 FINCH WEST, NaNs). Luckily there don't seem to be too many affected records, though the NaNs are not shown.

```

[ ]: delays_w_reasons['line'].value_counts()

[ ]: YU                8880
     BD                5734
     SHP               657
     SRT               656
     YU/BD             346
     YUS                18
     YU / BD           17
     YU & BD LINES      1

```

```

41 KEELE 1
52 1
35 JANE 1
999 1
YUS & BD 1
36 FINCH WEST 1
SHEP 1
YONGE/UNIVERSITY/BLOOR 1
YUS/BD 1
Name: line, dtype: int64

```

3.9.1 .loc[] and isna()

To find the records with no line, we can use `.loc[]`, which lets us access rows and columns with either a boolean array or row/column labels.

In this case, the boolean array is the product of the `isna()` Series method.

```
[ ]: # access rows of data where line is NA
delays_w_reasons.loc[delays_w_reasons['line'].isna()]
```

```
[ ]:
   date    time  day      station  code  min_delay  \
495  2021-01-13  15:22  Wednesday    FINCH WEST STATION  MUSAN      3
513  2021-01-13  22:08  Wednesday    EGLINTON WEST STATION  PUMEL      0
1044  2021-01-27  22:00  Wednesday  YONGE-UNIVERSITY AND B    MUO      0
1045  2021-01-27  23:00  Wednesday      FINCH STATION    MUO      0
1362  2021-02-04  01:45  Thursday    LAWRENCE STATION    TUSC      0
1679  2021-02-11  01:12  Thursday    GREENWOOD CARHOUSE  MUIE      0
2179  2021-02-22  08:27   Monday    BICHMOUNT DIVISION  MUIE      0
2204  2021-02-22  22:33   Monday      BLOOR STATION    SUAP      4
2206  2021-02-22  23:36   Monday    EGLINTON STATION    MUO      0
3039  2021-03-17  05:15  Wednesday    INGLIS BUILDING  PUMEL      0
3330  2021-03-24  19:13  Wednesday    INGLIS BUILDING  PUMEL      0
3407  2021-03-26  09:03   Friday  WILSON YARD (SOUTH TAI  PUTO      0
3557  2021-03-30  00:36   Tuesday    INGLIS BUILDING  PUMEL      0
3944  2021-04-08  23:45  Thursday    DAVISVILLE YARD  MUIE      0
4097  2021-04-13  10:57   Tuesday    SPADINA STATION    SUAE      0
4119  2021-04-13  22:00   Tuesday  YONGE-UNIVERSITY AND B    MUO      0
4336  2021-04-19  23:00   Monday  SHEPPARD WEST TO LAWRE  MUO      0
4748  2021-04-29  22:00  Thursday  YONGE UNIVERSITY SPADI  MUO      0
5312  2021-05-15  05:05  Saturday    SPADINA BD STATION  MUNCA      0
5448  2021-05-18  20:15   Tuesday    VAUGHAN MC STATION  MUWR      0
5484  2021-05-19  18:11  Wednesday    QUEEN'S QUAY STATION  PUMEL      0
5642  2021-05-23  23:19   Sunday      ST ANDREW STATION  SUDP      0
5685  2021-05-25  00:19   Tuesday    DUNDA WEST STATION    SUO      0
6042  2021-06-02  22:28  Wednesday      WARDEN STATION  MUIRS      0
6046  2021-06-02  00:56  Wednesday      BAY STATION    SUO      0

```

6540	2021-06-14	22:43	Monday	YONGE BD STATION	MUIS	0
6560	2021-06-15	07:15	Tuesday	SUBWAY OPERATIONS BUIL	PUMEL	0
7137	2021-06-28	01:03	Monday	COXWELL STATION	MUNCA	0
7766	2021-07-14	03:51	Wednesday	TRANSIT CONTROL CENTRE	PUSO	0
8889	2021-08-11	07:46	Wednesday	TRANSIT CONTROL	MUIE	0
9628	2021-08-29	15:49	Sunday	YORKDALE STATION	SUPOL	0
9629	2021-08-29	16:13	Sunday	YORK MILLS STATION	MUO	0
9780	2021-09-01	20:35	Wednesday	MAIN STREET AND UNION	MUO	0
9789	2021-09-01	22:14	Wednesday	UNION AND KENNEDY STAT	MUO	0
10336	2021-09-13	17:20	Monday	MCBRIEN BUILDING	SUO	0
10951	2021-09-26	15:50	Sunday	WILSON STATION	PUOPO	0
11223	2021-10-01	00:33	Friday	WELLESLEY STATION	SUDP	0
12533	2021-10-28	14:18	Thursday	VICTORIA PARK STATION	MUIS	0
12826	2021-11-02	12:22	Tuesday	GREENWOOD SHOP	MUIE	0
13007	2021-11-05	08:59	Friday	BLOOR STATION	MUIRS	0
13080	2021-11-06	18:41	Saturday	KENNEDY BD STATION	MUO	0
13273	2021-11-10	16:25	Wednesday	SUMMERHILL STATION	TUS	3
13402	2021-11-12	20:42	Friday	CLOSURES BUILDING	MUIE	0
13410	2021-11-12	00:02	Friday	TRANSIT CONTROL	MUO	0
14177	2021-11-25	21:14	Thursday	WILSON CARHOUSE	MUIE	0
14371	2021-11-29	05:10	Monday	GO PROTOCOL	MUO	0
14935	2021-12-08	06:00	Wednesday	TORONTO TRANSIT COMMIS	MUO	0
14952	2021-12-08	13:58	Wednesday	KIPLING STATION	MUIS	0
14967	2021-12-08	17:14	Wednesday	QUEEN'S PARK STATION	MUO	0
15581	2021-12-19	00:42	Sunday	WILSON CARHOUSE	MUO	0
15623	2021-12-20	16:23	Monday	YONGE-SHEPPARD (LINE 4	MUIRS	0
16332	2021-12-31	14:34	Friday	GO PROTOCOL	MUO	0

	min_gap	bound	line	vehicle \
495	6	S	NaN	5751
513	0	NaN	NaN	0
1044	0	NaN	NaN	0
1045	0	NaN	NaN	0
1362	0	S	NaN	5596
1679	0	NaN	NaN	0
2179	0	NaN	NaN	0
2204	9	N	NaN	6006
2206	0	NaN	NaN	0
3039	0	NaN	NaN	0
3330	0	NaN	NaN	0
3407	0	NaN	NaN	0
3557	0	NaN	NaN	0
3944	0	NaN	NaN	0
4097	0	NaN	NaN	0
4119	0	NaN	NaN	0
4336	0	NaN	NaN	0
4748	0	NaN	NaN	0

5312	0	NaN	NaN	0
5448	0	NaN	NaN	0
5484	0	NaN	NaN	0
5642	0	NaN	NaN	0
5685	0	E	NaN	0
6042	0	NaN	NaN	0
6046	0	NaN	NaN	0
6540	0	NaN	NaN	0
6560	0	NaN	NaN	0
7137	0	NaN	NaN	0
7766	0	NaN	NaN	0
8889	0	NaN	NaN	0
9628	0	NaN	NaN	0
9629	0	NaN	NaN	0
9780	0	NaN	NaN	0
9789	0	NaN	NaN	0
10336	0	NaN	NaN	0
10951	0	N	NaN	5471
11223	0	NaN	NaN	0
12533	0	NaN	NaN	0
12826	0	NaN	NaN	0
13007	0	S	NaN	0
13080	0	NaN	NaN	0
13273	6	N	NaN	6501
13402	0	NaN	NaN	0
13410	0	NaN	NaN	0
14177	0	NaN	NaN	0
14371	0	NaN	NaN	0
14935	0	NaN	NaN	0
14952	0	NaN	NaN	0
14967	0	NaN	NaN	0
15581	0	NaN	NaN	0
15623	0	NaN	NaN	0
16332	0	NaN	NaN	0

	code_description	sub_or_srt	year	\
495	Unsanitary Vehicle	SUB	2021	
513	Escalator/Elevator Incident	SUB	2021	
1044	Miscellaneous Other	SUB	2021	
1045	Miscellaneous Other	SUB	2021	
1362	Operator Overspeeding	SUB	2021	
1679	Injured Employee	SUB	2021	
2179	Injured Employee	SUB	2021	
2204	Assault / Patron Involved	SUB	2021	
2206	Miscellaneous Other	SUB	2021	
3039	Escalator/Elevator Incident	SUB	2021	
3330	Escalator/Elevator Incident	SUB	2021	

3407	T&S Other	SUB	2021
3557	Escalator/Elevator Incident	SUB	2021
3944	Injured Employee	SUB	2021
4097	Assault / Employee Involved	SUB	2021
4119	Miscellaneous Other	SUB	2021
4336	Miscellaneous Other	SUB	2021
4748	Miscellaneous Other	SUB	2021
5312	NaN	NaN	2021
5448	Work Refusal	SUB	2021
5484	Escalator/Elevator Incident	SUB	2021
5642	Disorderly Patron	SUB	2021
5685	Passenger Other	SUB	2021
6042	Injured or ill Customer (In Station) - Medical...	SUB	2021
6046	Passenger Other	SUB	2021
6540	Injured or ill Customer (In Station) - Transpo...	SUB	2021
6560	Escalator/Elevator Incident	SUB	2021
7137	NaN	NaN	2021
7766	S/E/C Department Other	SUB	2021
8889	Injured Employee	SUB	2021
9628	Held By Polce - Non-TTC Related	SUB	2021
9629	Miscellaneous Other	SUB	2021
9780	Miscellaneous Other	SUB	2021
9789	Miscellaneous Other	SUB	2021
10336	Passenger Other	SUB	2021
10951	OPTO (COMMS) Train Door Monitoring	SUB	2021
11223	Disorderly Patron	SUB	2021
12533	Injured or ill Customer (In Station) - Transpo...	SUB	2021
12826	Injured Employee	SUB	2021
13007	Injured or ill Customer (In Station) - Medical...	SUB	2021
13080	Miscellaneous Other	SUB	2021
13273	Crew Unable to Maintain Schedule	SUB	2021
13402	Injured Employee	SUB	2021
13410	Miscellaneous Other	SUB	2021
14177	Injured Employee	SUB	2021
14371	Miscellaneous Other	SUB	2021
14935	Miscellaneous Other	SUB	2021
14952	Injured or ill Customer (In Station) - Transpo...	SUB	2021
14967	Miscellaneous Other	SUB	2021
15581	Miscellaneous Other	SUB	2021
15623	Injured or ill Customer (In Station) - Medical...	SUB	2021
16332	Miscellaneous Other	SUB	2021

	hour_delay	month	hour
495	0.05	1	15
513	0.00	1	22
1044	0.00	1	22
1045	0.00	1	23

1362	0.00	2	1
1679	0.00	2	1
2179	0.00	2	8
2204	0.07	2	22
2206	0.00	2	23
3039	0.00	3	5
3330	0.00	3	19
3407	0.00	3	9
3557	0.00	3	0
3944	0.00	4	23
4097	0.00	4	10
4119	0.00	4	22
4336	0.00	4	23
4748	0.00	4	22
5312	0.00	5	5
5448	0.00	5	20
5484	0.00	5	18
5642	0.00	5	23
5685	0.00	5	0
6042	0.00	6	22
6046	0.00	6	0
6540	0.00	6	22
6560	0.00	6	7
7137	0.00	6	1
7766	0.00	7	3
8889	0.00	8	7
9628	0.00	8	15
9629	0.00	8	16
9780	0.00	9	20
9789	0.00	9	22
10336	0.00	9	17
10951	0.00	9	15
11223	0.00	10	0
12533	0.00	10	14
12826	0.00	11	12
13007	0.00	11	8
13080	0.00	11	18
13273	0.05	11	16
13402	0.00	11	20
13410	0.00	11	0
14177	0.00	11	21
14371	0.00	11	5
14935	0.00	12	6
14952	0.00	12	13
14967	0.00	12	17
15581	0.00	12	0
15623	0.00	12	16


```
16332          0.00      12      14
```

`.loc[]` also lets us access data by label, with row conditions first and column conditions second.

```
[ ]: (delays_w_reasons.loc[delays_w_reasons['line'].isna(), # filter rows
                        ['time', 'station', 'line']] # get columns
      .head()) # first 5 lines to save space
```

```
[ ]:      time      station line
495   15:22   FINCH WEST STATION  NaN
513   22:08  EGLINTON WEST STATION  NaN
1044  22:00  YONGE-UNIVERSITY AND B  NaN
1045  23:00      FINCH STATION  NaN
1362  01:45   LAWRENCE STATION  NaN
```

3.9.2 query()

Alternatively, we can use the DataFrame `query()` method, which takes a filter condition as a string, and returns a DataFrame of records that met the condition. `query()` is slower than `loc[]`, but it can be easier to read.

```
[ ]: delays_w_reasons['line'].unique()
```

```
[ ]: array(['YU', 'BD', 'SHP', 'SRT', 'YU/BD', nan, 'YONGE/UNIVERSITY/BLOOR',
          'YU / BD', 'YUS', '999', 'SHEP', '36 FINCH WEST', 'YUS & BD',
          'YU & BD LINES', '35 JANE', '52', '41 KEELE', 'YUS/BD'],
          dtype=object)
```

```
[ ]: delays_w_reasons['line'].isna()
```

```
[ ]: 0      False
      1      False
      2      False
      3      False
      4      False
      ...
16365  False
16366  False
16367  False
16368  False
16369  False
Name: line, Length: 16370, dtype: bool
```

```
[ ]: # slower than .loc, but can be easier to read
delays_w_reasons.query('line.isna()', engine='python').head()
```

```
[ ]:      date    time    day      station    code  min_delay  \
495  2021-01-13  15:22  Wednesday    FINCH WEST STATION  MUSAN          3
513  2021-01-13  22:08  Wednesday    EGLINTON WEST STATION  PUMEL          0
1044 2021-01-27  22:00  Wednesday  YONGE-UNIVERSITY AND B    MUO          0
1045 2021-01-27  23:00  Wednesday    FINCH STATION    MUO          0
1362 2021-02-04  01:45   Thursday    LAWRENCE STATION   TUSC          0

      min_gap bound line  vehicle      code_description sub_or_srt  \
495         6      S  NaN    5751      Unsanitary Vehicle      SUB
513         0     NaN  NaN      0  Escalator/Elevator Incident      SUB
1044         0     NaN  NaN      0      Miscellaneous Other      SUB
1045         0     NaN  NaN      0      Miscellaneous Other      SUB
1362         0      S  NaN    5596      Operator Overspeeding      SUB

      year  hour_delay  month  hour
495   2021         0.05      1    15
513   2021         0.00      1    22
1044  2021         0.00      1    22
1045  2021         0.00      1    23
1362  2021         0.00      2     1
```

3.9.3 dropna()

In this case, the number of records without lines is relatively small. Most do not have delay durations. Some appear to be at rail yards, i.e. not on a rail line. For our analysis, we may drop them with the `dropna()` DataFrame method. We can drop rows missing lines by passing a `subset`.

```
[ ]: delays_w_reasons = delays_w_reasons.dropna(subset=['line'])
```

3.9.4 Filtering data with `.loc[]` and `isin()`

We can use `.loc[]` to create a `delays` DataFrame without the invalid lines. To do this, we first create a list of values to exclude, then pass the list to the Series `isin()` method. Finally, we negate the expression, and assign the output back to `delays_w_reasons`.

Note: The negation operator here is `~`, not `!`. The `and` and `or` operators are different as well: `&` and `|` respectively.

```
[ ]: # set up filter list
filter_list = ['999', '36 FINCH WEST', '35 JANE', '52', '41 KEELE']

[ ]: # filter out records with invalid lines
delays_w_reasons = delays_w_reasons.loc[~delays_w_reasons['line'].
    ↪isin(filter_list)]
delays_w_reasons['line'].unique()
```

```
[ ]: array(['YU', 'BD', 'SHP', 'SRT', 'YU/BD', 'YONGE/UNIVERSITY/BLOOR',
          'YU / BD', 'YUS', 'SHEP', 'YUS & BD', 'YU & BD LINES', 'YUS/BD'],
          dtype=object)
```

3.9.5 Replacing values with `str.replace()`

To standardize the YU/BD values, we can replace the less common ones. One way to do this is by selecting the line Series and using `str.replace()`, like below for "YUS".

```
[ ]: delays_w_reasons['line'] = (delays_w_reasons['line']
                                .str.replace('YUS', 'YU'))
delays_w_reasons['line'].unique()
```

```
[ ]: array(['YU', 'BD', 'SHP', 'SRT', 'YU/BD', 'YONGE/UNIVERSITY/BLOOR',
          'YU / BD', 'SHEP', 'YU & BD', 'YU & BD LINES'], dtype=object)
```

Another is to assign "YU/BD" to values selected by `.loc[]`

```
[ ]: yubd_list = ['YONGE/UNIVERSITY/BLOOR',
                 'YU / BD',
                 'YU & BD',
                 'YU & BD LINES']

# check the .loc[] selection
delays_w_reasons.loc[delays_w_reasons['line'].isin(yubd_list), 'line']
```

```
[ ]: 590      YONGE/UNIVERSITY/BLOOR
      852      YU / BD
      1137     YU / BD
      1628     YU / BD
      1672     YU / BD
      1700     YU / BD
      6725     YU / BD
      7469     YU / BD
      8034     YU & BD
      8301     YU / BD
      8341     YU / BD
      8463     YU / BD
      9164     YU / BD
      9541     YU & BD LINES
      9839     YU / BD
     10792     YU / BD
     11119     YU / BD
     11299     YU / BD
     12128     YU / BD
     15574     YU / BD
Name: line, dtype: object
```

```
[ ]: delays_w_reasons.loc[delays_w_reasons['line'].isin(yubd_list), 'line'] = 'YU/BD'
delays_w_reasons['line'].unique()
```

```
[ ]: array(['YU', 'BD', 'SHP', 'SRT', 'YU/BD', 'SHEP'], dtype=object)
```

3.10 Grouping

A core workflow in `pandas` is *split-apply-combine*: * **splitting** data into groups * **applying** a function to each group, such as calculating group sums, standardizing data, or filtering out some groups * **combining** the results into a data structure

This workflow starts by grouping data by calling the `groupby()` method. We'll pass in a column name or list of names to group by.

```
[ ]: line_groups = delays_w_reasons.groupby('line')
```

`groupby()` returns a grouped `DataFrame` that we can use to calculate groupwise statistics. The grouping column values become indexes, or row labels. **Note that this grouped `DataFrame` still references the original, so mutating one affects the other.**

```
[ ]: # how many hours of delays did each line have in 2021?
line_groups['hour_delay'].sum()
```

```
[ ]: line
BD      329.47
SHEP     0.00
SHP     28.43
SRT     57.82
YU     477.50
YU/BD    0.00
Name: hour_delay, dtype: float64
```

We can group by more than one column by passing a list into `groupby()`. Data is grouped in the order of column names.

```
[ ]: # group by line first, then reason code description
line_code_groups = delays_w_reasons.groupby(['line', 'code_description'])
```

3.10.1 Chaining methods and `unstack()`ing

We can *chain* methods together for convenience and code readability. Here, we calculate the `size()` of each group, then `unstack()` the resulting Series by the first part of the row label, line. The `tail()` method is added to the end so that the output takes less screen space.

```
[ ]: # view the number of delays by reason and line
line_code_groups.size().unstack(0).tail()
```

```
[ ]: line
code_description      BD  SHEP  SHP  SRT    YU  YU/BD
Work Refusal          4.0   NaN  1.0  NaN  12.0   NaN
Work Vehicle          3.0   NaN  NaN  NaN   7.0   NaN
Work Zone Problems - Signals  4.0   NaN  4.0  NaN   5.0   NaN
Work Zone Problems - Track  12.0   NaN  NaN  NaN  29.0   NaN
Yard/Carhouse Related Problems 17.0   NaN  NaN  NaN  15.0   NaN
```

3.10.2 agg()regating

So far, we have applied one function at a time. The `agg()` DataFrame method lets us apply multiple functions on different columns at once.

`agg()`'s argument syntax is a little unusual. It follows this pattern:

```
DataFrame.agg(agg_colname=('column_to_aggregate', 'aggregation_function_name'),
              agg_colname2=('col_to_agg2', 'agg_func_name'))
```

```
[ ]: delay_summary = (delays_w_reasons
                      .groupby('date')
                      .agg(delay_count=('station', 'count'),
                          total_delay_min=('min_delay', 'sum'),
                          mean_delay_min=('min_delay', 'mean')))

delay_summary.head()
```


```
[ ]:
      delay_count  total_delay_min  mean_delay_min
date
2021-01-01        36             159         4.416667
2021-01-02        49             284         5.795918
2021-01-03        19              51         2.684211
2021-01-04        41             284         6.926829
2021-01-05        40             298         7.450000
```

3.11 Plotting

The summary table we just generated can be easily plotted within `pandas`. Since the index contains dates, `pandas` automatically knows to plot values as time series data, with the dates in the x-axis -- we just have to call the `plot()` method.

```
[ ]: delay_summary.plot()
```


```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe3fcb623d0>
```



To create a separate plot for each column, we can specify `subplots=True`

```
[ ]: delay_summary.plot(subplots=True)
```

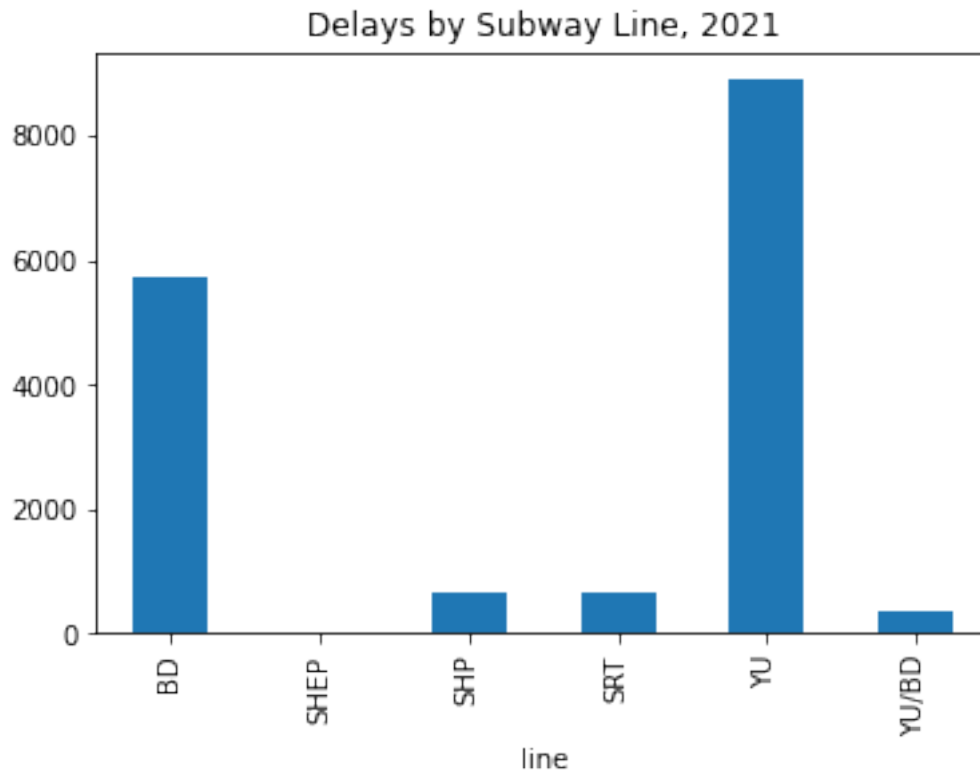
```
[ ]: array([<matplotlib.axes._subplots.AxesSubplot object at 0x7fe3fca515d0>,
          <matplotlib.axes._subplots.AxesSubplot object at 0x7fe3fc52db50>,
          <matplotlib.axes._subplots.AxesSubplot object at 0x7fe3fca1bf90>],
          dtype=object)
```



We can plot other aggregations too. Below, we use `line_groups` and calculate the size of each group, i.e., the number of delays reported on each line. Then we plot the data, telling `pandas` that the plot kind should be a bar graph, with TTC lines should in the x-axis. We also pass in a title.

```
[ ]: (line_groups
      .size()
      .plot(x='line',
            kind='bar',
            title='Delays by Subway Line, 2021'))
```


```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe3fc287ed0>
```



It is possible to sum up and plot the total delay time, in hours, by line.

```
[ ]: (delays_w_reasons
      .groupby('line')['min_delay']
      .sum()
      .plot(x='line', kind='bar'))
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe3fc198350>
```

4 Writing to file

4.1 Exporting DataFrames

DataFrames have `to_[file format]()` methods, analogous to `pandas` read functions. The counterpart to `pd.read_csv()`, for instance, is `DataFrame.to_csv()`. The export methods generally take a file path to save to as their first argument. Additional arguments vary a bit by export format, but `index` is a common useful one. It takes a boolean of whether to write the index to file -- set it to `False` if the index is the numbered default.

Two additional useful parameters in `DataFrame.to_csv()` and `DataFrame.to_excel()` are `na_rep`, which takes a string to use for null values, and `columns`, which lets us write out a subset of columns.

```
[ ]: # write delays_w_reasons to an Excel file
delays_w_reasons.to_excel('/content/data/ttc_subway_delays_w_reasons.xlsx',
    ↪index=False)
```

5 More wrangling

5.1 Neighbourhood Profiles

The bike theft data includes neighbourhood identifiers. These neighbourhoods are designated by City of Toronto, which publishes neighbourhood demographic profiles. Let's get this data to start investigating if neighbourhoods with more bike theft reports simply have higher populations. In the process, we will reinforce what we learned so far. We will also learn about two last ways to reshape data: `melt()`, which rearranges data from a wide format to a long format; and `pivot()`, which reorganizes data based on index and column values.

5.2 Getting data

Let's load the neighbourhood data and explore it.

```
[ ]: profiles = pd.read_csv('/content/data/neighbourhood-profiles-2016-140-model.  
    ↪ csv')
```

```
[ ]: profiles.shape
```

```
[ ]: (2383, 146)
```

The neighbourhood profiles are in an unusual format. Neighbourhood names are in the columns, while attribute fields are rows, and there are thousands of attributes.

```
[ ]: profiles.head()
```

```
[ ]: 

|   | _id | Category                  | Topic \                   |
|---|-----|---------------------------|---------------------------|
| 0 | 1   | Neighbourhood Information | Neighbourhood Information |
| 1 | 2   | Neighbourhood Information | Neighbourhood Information |
| 2 | 3   | Population                | Population and dwellings  |
| 3 | 4   | Population                | Population and dwellings  |
| 4 | 5   | Population                | Population and dwellings  |


|   | Data Source                    | Characteristic \            |
|---|--------------------------------|-----------------------------|
| 0 | City of Toronto                | Neighbourhood Number        |
| 1 | City of Toronto                | TSNS2020 Designation        |
| 2 | Census Profile 98-316-X2016001 | Population, 2016            |
| 3 | Census Profile 98-316-X2016001 | Population, 2011            |
| 4 | Census Profile 98-316-X2016001 | Population Change 2011-2016 |


|   | City of Toronto | Agincourt North | Agincourt South-Malvern West \ |
|---|-----------------|-----------------|--------------------------------|
| 0 | NaN             | 129             | 128                            |
| 1 | NaN             | No Designation  | No Designation                 |
| 2 | 2,731,571       | 29,113          | 23,757                         |
| 3 | 2,615,060       | 30,279          | 21,988                         |
| 4 | 4.50%           | -3.90%          | 8.00%                          |


```

	Alderwood	Annex Banbury-Don Mills	Bathurst Manor	\
0	20	95	42	34
1	No Designation	No Designation	No Designation	No Designation
2	12,054	30,526	27,695	15,873
3	11,904	29,177	26,918	15,434
4	1.30%	4.60%	2.90%	2.80%

	Bay Street Corridor	Bayview Village	Bayview Woods-Steeles	\
0	76	52	49	
1	No Designation	No Designation	No Designation	
2	25,797	21,396	13,154	
3	19,348	17,671	13,530	
4	33.30%	21.10%	-2.80%	

	Bedford Park-Nortown	Beechborough-Greenbrook	Bendale	\
0	39	112	127	
1	No Designation	NIA	No Designation	
2	23,236	6,577	29,960	
3	23,185	6,488	27,876	
4	0.20%	1.40%	7.50%	

	Birchcliffe-Cliffside	Black Creek	Blake-Jones	Briar Hill-Belgravia	\
0	122	24	69	108	
1	No Designation	NIA	No Designation	No Designation	
2	22,291	21,737	7,727	14,257	
3	21,856	22,057	7,763	14,302	
4	2.00%	-1.50%	-0.50%	-0.30%	

	Bridle Path-Sunnybrook-York Mills	Broadview North	Brookhaven-Amesbury	\
0	41	57	30	
1	No Designation	No Designation	No Designation	
2	9,266	11,499	17,757	
3	8,713	11,563	17,787	
4	6.30%	-0.60%	-0.20%	

	Cabbagetown-South St. James Town	Caledonia-Fairbank	Casa Loma	\
0	71	109	96	
1	No Designation	No Designation	No Designation	
2	11,669	9,955	10,968	
3	12,053	9,851	10,487	
4	-3.20%	1.10%	4.60%	

	Centennial Scarborough Church-Yonge Corridor	Clairlea-Birchmount	\
0	133	75	120
1	No Designation	No Designation	No Designation
2	13,362	31,340	26,984

3	13,093	28,349	24,770
4	2.10%	10.60%	8.90%

	Clanton Park	Cliffcrest Corso Italia-Davenport	Danforth \
0	33	123	92 66
1	No Designation	No Designation	No Designation No Designation
2	16,472	15,935	14,133 9,666
3	14,612	15,703	13,743 9,444
4	12.70%	1.50%	2.80% 2.40%

	Danforth East York Don Valley Village	Dorset Park \
0	59 47	126
1	No Designation No Designation	Emerging Neighbourhood
2	17,180 27,051	25,003
3	16,712 26,739	24,363
4	2.80% 1.20%	2.60%

	Dovercourt-Wallace Emerson-Junction Downsview-Roding-CFB	Dufferin Grove \
0	93 26	83
1	No Designation	NIA No Designation
2	36,625 35,052	11,785
3	34,631 34,659	11,449
4	5.80% 1.10%	2.90%

	East End-Danforth Edenbridge-Humber Valley Eglinton East Elms-Old Rexdale \
0	62 9 138 5
1	No Designation No Designation NIA NIA
2	21,381 15,535 22,776 9,456
3	20,839 14,943 22,829 9,550
4	2.60% 4.00% -0.20% -1.00%

	Englemount-Lawrence Eringate-Centennial-West Deane Etobicoke West Mall \
0	32 11 13
1	Emerging Neighbourhood No Designation No Designation
2	22,372 18,588 11,848
3	22,086 18,810 10,927
4	1.30% -1.20% 8.40%

	Flemingdon Park Forest Hill North Forest Hill South Glenfield-Jane Heights \
0	44 102 101 25
1	NIA No Designation No Designation NIA
2	21,933 12,806 10,732 30,491
3	22,168 12,474 10,926 31,390
4	-1.10% 2.70% -1.80% -2.90%

	Greenwood-Coxwell Guildwood Henry Farm High Park North \
0	65 140 53 88

1	No Designation	No Designation	No Designation	No Designation
2	14,417	9,917	15,723	22,162
3	14,083	9,816	11,333	21,292
4	2.40%	1.00%	38.70%	4.10%

	High Park-Swansea	Highland Creek	Hillcrest Village	\
0	87	134	48	
1	No Designation	No Designation	No Designation	
2	23,925	12,494	16,934	
3	21,740	13,097	17,656	
4	10.10%	-4.60%	-4.10%	

	Humber Heights-Westmount	Humber Summit	Humbermede	Humewood-Cedarvale	\
0	8	21	22	106	
1	Emerging Neighbourhood	NIA	NIA	No Designation	
2	10,948	12,416	15,545	14,365	
3	10,583	12,525	15,853	14,108	
4	3.40%	-0.90%	-1.90%	1.80%	

	Ionview	Islington-City Centre West	Junction Area	Keelesdale-Eglinton West	\
0	125	14	90	110	
1	NIA	No Designation	No Designation	NIA	
2	13,641	43,965	14,366	11,058	
3	13,091	38,084	14,027	10,638	
4	4.20%	15.40%	2.40%	3.90%	

	Kennedy Park	Kensington-Chinatown	Kingsview Village-The Westway	\
0	124	78	6	
1	NIA	No Designation	NIA	
2	17,123	17,945	22,000	
3	17,058	18,495	21,723	
4	0.40%	-3.00%	1.30%	

	Kingsway South	Lambton Baby Point	L'Amoreaux	Lansing-Westgate	\
0	15	114	117	38	
1	No Designation	No Designation	Emerging Neighbourhood	No Designation	
2	9,271	7,985	43,993	16,164	
3	9,170	7,921	44,919	14,642	
4	1.10%	0.80%	-2.10%	10.40%	

	Lawrence Park North	Lawrence Park South	Leaside-Bennington	Little Portugal	\
0	105	103	56	84	
1	No Designation	No Designation	No Designation	No Designation	
2	14,607	15,179	16,828	15,559	
3	14,541	15,070	17,011	12,050	
4	0.50%	0.70%	-1.10%	29.10%	

	Long Branch	Malvern	Maple Leaf	Markland Wood \
0	19	132	29	12
1	No Designation	Emerging Neighbourhood	No Designation	No Designation
2	10,084	43,794	10,111	10,554
3	9,632	45,086	10,197	10,436
4	4.70%	-2.90%	-0.80%	1.10%

	Milliken Mimico (includes Humber Bay Shores)	Morningside \
0	130	17 135
1	No Designation	No Designation NIA
2	26,572	33,964 17,455
3	27,167	26,541 17,587
4	-2.20%	28.00% -0.80%

	Moss Park	Mount Dennis	Mount Olive-Silverstone-Jamestown \
0	73	115	2
1	No Designation	NIA	NIA
2	20,506	13,593	32,954
3	16,306	13,145	32,788
4	25.80%	3.40%	0.50%

	Mount Pleasant East	Mount Pleasant West	New Toronto	Newtonbrook East \
0	99	104	18	50
1	No Designation	No Designation	No Designation	No Designation
2	16,775	29,658	11,463	16,097
3	15,982	28,593	10,900	16,423
4	5.00%	3.70%	5.20%	-2.00%

	Newtonbrook West	Niagara North	Riverdale	North St. James Town \
0	36	82	68	74
1	No Designation	No Designation	No Designation	No Designation
2	23,831	31,180	11,916	18,615
3	23,052	21,274	12,191	17,832
4	3.40%	46.60%	-2.30%	4.40%

	Oakridge	Oakwood Village	O'Connor-Parkview	Old East York \
0	121	107	54	58
1	NIA	No Designation	No Designation	No Designation
2	13,845	21,210	18,675	9,233
3	13,497	21,073	18,316	9,118
4	2.60%	0.70%	2.00%	1.30%

	Palmerston-Little Italy	Parkwoods-Donalda	Pelmo Park-Humberlea \
0	80	45	23
1	No Designation	No Designation	No Designation
2	13,826	34,805	10,722
3	13,746	34,617	8,710

4	0.60%	0.50%	23.10%	
	Playter Estates-Danforth	Pleasant View	Princess-Rosethorn	Regent Park \
0	67	46	10	72
1	No Designation	No Designation	No Designation	NIA
2	7,804	15,818	11,051	10,803
3	7,653	16,144	11,197	10,007
4	2.00%	-2.00%	-1.30%	8.00%
	Rexdale-Kipling	Rockcliffe-Smythe	Roncesvalles	Rosedale-Moore Park \
0	4	111	86	98
1	No Designation	NIA	No Designation	No Designation
2	10,529	22,246	14,974	20,923
3	10,488	22,267	15,050	20,631
4	0.40%	-0.10%	-0.50%	1.40%
	Rouge	Runnymede-Bloor	West Village	Rustic Scarborough Village \
0	131	89	28	139
1	No Designation	No Designation	NIA	NIA
2	46,496	10,070	9,941	16,724
3	45,912	9,632	9,951	16,609
4	1.30%	4.50%	-0.10%	0.70%
	South Parkdale	South Riverdale	St. Andrew-Windfields	Steeles \
0	85	70	40	116
1	NIA	No Designation	No Designation	Emerging Neighbourhood
2	21,849	27,876	17,812	24,623
3	21,251	25,642	17,958	25,017
4	2.80%	8.70%	-0.80%	-1.60%
	Stonegate-Queensway	Tam O'Shanter-Sullivan	Taylor-Massey	The Beaches \
0	16	118	61	63
1	No Designation	No Designation	NIA	No Designation
2	25,051	27,446	15,683	21,567
3	24,691	27,398	15,594	21,130
4	1.50%	0.20%	0.60%	2.10%
	Thistletown-Beaumont Heights	Thorncliffe Park	Trinity-Bellwoods	\
0	3	55	81	
1	NIA	NIA	No Designation	
2	10,360	21,108	16,556	
3	10,138	19,225	16,802	
4	2.20%	9.80%	-1.50%	
	University	Victoria Village	Waterfront Communities-The Island	\
0	79	43	77	
1	No Designation	NIA	No Designation	

2	7,607	17,510	65,913
3	7,782	17,182	43,361
4	-2.20%	1.90%	52.00%

	West Hill	West Humber-Clairville	Westminster-Branson	Weston \
0	136	1	35	113
1	NIA	No Designation	Emerging Neighbourhood	NIA
2	27,392	33,312	26,274	17,992
3	26,547	34,100	25,446	18,170
4	3.20%	-2.30%	3.30%	-1.00%

	Weston-Pelham Park	Wexford/Maryvale	Willowdale East	Willowdale West \
0	91	119	51	37
1	NIA	No Designation	No Designation	No Designation
2	11,098	27,917	50,434	16,936
3	12,010	27,018	45,041	15,004
4	-7.60%	3.30%	12.00%	12.90%

	Willowridge-Martingrove-Richview	Woburn	Woodbine Corridor	Woodbine-Lumsden \
0	7	137	64	60
1	No Designation	NIA	No Designation	No Designation
2	22,156	53,485	12,541	7,865
3	21,343	53,350	11,703	7,826
4	3.80%	0.30%	7.20%	0.50%

	Wychwood	Yonge-Eglinton	Yonge-St.Clair	York University Heights \
0	94	100	97	27
1	No Designation	No Designation	No Designation	NIA
2	14,349	11,817	12,528	27,593
3	13,986	10,578	11,652	27,713
4	2.60%	11.70%	7.50%	-0.40%

	Yorkdale-Glen Park
0	31
1	Emerging Neighbourhood
2	14,804
3	14,687
4	0.80%

Because of the layout and formatting characters, all of the numeric values have been read in as text data. It also looks like the characteristics are not unique.

```
[ ]: profiles.dtypes.value_counts()
```

```
[ ]: object    145
      int64     1
      dtype: int64
```



```
[ ]: len(profiles['Characteristic'].unique())
```

```
[ ]: 1651
```

5.3 Removing extra whitespace

The characteristic values contain extra whitespace. Let's remove the whitespace up with `str.strip()`.

```
[ ]: # the whitespace is easier to see in a list than a Series  
list(profiles['Characteristic'][95:100])
```

```
[ ]: ['    Female parent',  
      '    Male parent',  
      'Couple census families in private households',  
      '    Couples with children',  
      '    1 child']
```

```
[ ]: profiles['Characteristic'] = profiles['Characteristic'].str.strip()  
  
# get the first 10 characteristics  
list(profiles['Characteristic'][95:100])
```

```
[ ]: ['Female parent',  
      'Male parent',  
      'Couple census families in private households',  
      'Couples with children',  
      '1 child']
```

5.4 Subsetting data

1651 characteristics is still a lot. Let's check out the categories to understand the areas covered.

```
[ ]: profiles['Category'].unique()
```

```
[ ]: array(['Neighbourhood Information', 'Population',  
          'Families, households and marital status', 'Language', 'Income',  
          'Immigration and citizenship', 'Visible minority', 'Ethnic origin',  
          'Aboriginal peoples', 'Education', 'Housing', 'Language of work',  
          'Labour', 'Journey to work', 'Mobility'], dtype=object)
```

"Journey to work" sounds relevant. We can use `.loc[]` to get the rows in that category, then select the Topic column and get its unique values.

```
[ ]: profiles.loc[profiles['Category'] == 'Journey to work']['Topic'].unique()
```

```
[ ]: array(['Commuting destination', 'Main mode of commuting',
          'Commuting duration', 'Time leaving for work'], dtype=object)
```

The "Population and dwellings" topic we saw in the DataFrame head looked promising as well. Let's check out the Characteristics in that topic.

```
[ ]: profiles.loc[profiles['Topic'] == 'Population and dwellings']['Characteristic'].
      ↪unique()
```

```
[ ]: array(['Population, 2016', 'Population, 2011',
          'Population Change 2011-2016', 'Total private dwellings',
          'Private dwellings occupied by usual residents',
          'Population density per square kilometre',
          'Land area in square kilometres'], dtype=object)
```

Now that we know what topics we're interested in, let's create a subset DataFrame limited to them. We'll use the `copy()` DataFrame method to leave the original data untouched.

```
[ ]: topics = ['Neighbourhood Information', 'Population and dwellings', 'Main mode
      ↪of commuting']

      # make sure it's an independent copy
      profiles_subset = profiles.copy()

      # get just the topics we're interested in
      profiles_subset = profiles_subset.loc[profiles['Topic'].isin(topics)]
      profiles_subset.shape
```

```
[ ]: (16, 146)
```

5.5 Reshaping data with `melt()`

Now we're ready to reshape our data. We can `drop()` the ID, data source, and category columns now.

To `melt()` a DataFrame, we specify `id_vars` -- the columns to keep as identifiers. All other columns are 'melted' into a new `variable` column. The values at `DataFrame[id_vars, variable_col]` move into a `value` column. We can change the names of the variable and value columns with the `var_name` and `value_name` arguments.

The pandas documentation provides an [illustrative example](#).

Let's `melt()` the profiles subset. We'll keep `Topic` and `Characteristic` as our `id_vars`. This will melt the neighbourhood names into the `variable` column, which we'll rename `Neighbourhood`.

```
[ ]: profiles_melt = (profiles_subset
      .drop(columns=['_id', 'Data Source', 'Category'])
      .melt(id_vars=['Topic', 'Characteristic'],
```

```
var_name='Neighbourhood'))
profiles_melt.head()
```

```
[ ]:      Topic      Characteristic  Neighbourhood \
0  Neighbourhood Information  Neighbourhood Number  City of Toronto
1  Neighbourhood Information  TSNS2020 Designation  City of Toronto
2  Population and dwellings    Population, 2016    City of Toronto
3  Population and dwellings    Population, 2011  City of Toronto
4  Population and dwellings  Population Change 2011-2016  City of Toronto

      value
0      NaN
1      NaN
2  2,731,571
3  2,615,060
4      4.50%
```

5.6 Reshaping data with pivot()

The profile data is looking much closer to what we want! The next step is to make the Topic/Characteristic the column header, `pivot()`ing the values. To do this, we specify the column(s) to use as the `index`, or row labels; the column(s) whose values we should use as `column` names, and which column our `values` come from.

Pivoting on two columns creates a multi-level column header, so we then drop the top Topic level with `droplevel()`. Finally, we `reset_index()` to make neighbourhood names a regular column.

```
[ ]: neighbourhoods = (profiles_melt.pivot(index='Neighbourhood',
      columns=['Topic', 'Characteristic'],
      values='value')
      .droplevel(0, axis=1) # remove topic col header
      .reset_index()) # make Neighbourhood a regular
↳ column
neighbourhoods.head()
```

```
[ ]: Characteristic      Neighbourhood Neighbourhood Number \
0      Agincourt North      129
1  Agincourt South-Malvern West  128
2      Alderwood      20
3      Annex      95
4  Banbury-Don Mills      42

Characteristic TSNS2020 Designation Population, 2016 Population, 2011 \
0      No Designation      29,113      30,279
1      No Designation      23,757      21,988
2      No Designation      12,054      11,904
3      No Designation      30,526      29,177
```

4	No Designation	27,695	26,918
---	----------------	--------	--------

Characteristic Population Change 2011-2016 Total private dwellings \

0	-3.90%	9,371
1	8.00%	8,535
2	1.30%	4,732
3	4.60%	18,109
4	2.90%	12,473

Characteristic Private dwellings occupied by usual residents \

0	9,120
1	8,136
2	4,616
3	15,934
4	12,124

Characteristic Population density per square kilometre \

0	3,929
1	3,034
2	2,435
3	10,863
4	2,775

Characteristic Land area in square kilometres \

0	7.41
1	7.83
2	4.95
3	2.81
4	9.98

Characteristic Total - Main mode of commuting for the employed labour force aged 15 years and over in private households with a usual place of work or no fixed workplace address - 25% sample data \

0	11,820
1	10,160
2	6,045
3	14,910
4	11,395

Characteristic Car, truck, van - as a driver Car, truck, van - as a passenger \

0	7,155	930
1	6,135	665
2	4,090	355
3	3,290	290
4	7,150	500

Characteristic Public transit Walked Bicycle Other method

0	3,350	265	70	45
1	2,985	280	35	65
2	1,285	195	65	65
3	6,200	3,200	1,675	225
4	2,945	615	65	140

5.7 Renaming all columns

Much better! These column names could be shorter, though. Let's rename them to be easier to work with. We could use the `rename()` DataFrame method, passing in a dictionary of old and new names. Since there isn't an easy renaming function, and some of the current names are very long, we will instead reassign a list of new names to the `columns` attribute of our DataFrame.

```
[ ]: # rename all columns
neighbourhoods.columns = ['neighbourhood',
                           'n_id',
                           'designation',
                           'pop_2016',
                           'pop_2011',
                           'pop_change',
                           'private_dwellings',
                           'occupied_dwellings',
                           'pop_dens',
                           'area',
                           'total_commuters',
                           'drive',
                           'car_passenger',
                           'transit',
                           'walk',
                           'bike',
                           'other']

neighbourhoods.columns
```

```
[ ]: Index(['neighbourhood', 'n_id', 'designation', 'pop_2016', 'pop_2011',
            'pop_change', 'private_dwellings', 'occupied_dwellings', 'pop_dens',
            'area', 'total_commuters', 'drive', 'car_passenger', 'transit', 'walk',
            'bike', 'other'],
           dtype='object')
```

5.8 Replacing values in multiple columns

All of the values in our neighbourhood data are text right now. Part of the problem is that numbers contain characters like commas and percentage signs. We can remove these from everywhere in our data with the DataFrame `replace()` method, which takes a string to look for and a replacement string. Normally, `replace()` looks for a perfect, full-string match. Since we're only looking for a substring match, we set `regex=True`.

```
[ ]: # for those comfortable with regex, ',|%' and '[,%]' also work
neighbourhoods = (neighbourhoods.replace(',', ' ', regex=True)
                  .replace('%', ' ', regex=True))
neighbourhoods.head(2)
```

```
[ ]:
      neighbourhood n_id      designation pop_2016 pop_2011 \
0      Agincourt North  129  No Designation   29113   30279
1  Agincourt South-Malvern West  128  No Designation   23757   21988

      pop_change private_dwellings occupied_dwellings pop_dens  area \
0      -3.90           9371           9120      3929  7.41
1       8.00           8535           8136      3034  7.83

      total_commuters drive car_passenger transit walk bike other
0           11820  7155           930   3350  265   70   45
1           10160  6135           665   2985  280   35   65
```

5.9 apply()ing a function to multiple columns

Now the numbers look like numbers, but they are still strings. We can convert them with `pd.to_numeric()`, which takes a Series and returns it as the most appropriate numeric data type. Doing this for columns one-by-one would be tedious. Instead, we can use the `apply()` DataFrame method to run a function on every column in a DataFrame. `apply()` takes the name of the function to apply and any arguments needed to run that function. We only want to convert from `pop_2016` onwards, so we'll use `.loc[]` to select the correct columns.

```
[ ]: # select all rows, columns from pop_2016 to end
neighbourhoods.loc[:, 'pop_2016':] = neighbourhoods.loc[:, 'pop_2016':].
    ↪ apply(pd.to_numeric)
neighbourhoods.head()
```

```
[ ]:
      neighbourhood n_id      designation pop_2016 pop_2011 \
0      Agincourt North  129  No Designation  29113.0  30279.0
1  Agincourt South-Malvern West  128  No Designation  23757.0  21988.0
2      Alderwood      20  No Designation  12054.0  11904.0
3      Annex         95  No Designation  30526.0  29177.0
4  Banbury-Don Mills   42  No Designation  27695.0  26918.0

      pop_change private_dwellings occupied_dwellings pop_dens  area \
0      -3.9           9371.0           9120.0   3929.0  7.41
1       8.0           8535.0           8136.0   3034.0  7.83
2       1.3           4732.0           4616.0   2435.0  4.95
3       4.6          18109.0          15934.0  10863.0  2.81
4       2.9          12473.0          12124.0   2775.0  9.98

      total_commuters drive car_passenger transit walk bike other
```

0	11820.0	7155.0	930.0	3350.0	265.0	70.0	45.0
1	10160.0	6135.0	665.0	2985.0	280.0	35.0	65.0
2	6045.0	4090.0	355.0	1285.0	195.0	65.0	65.0
3	14910.0	3290.0	290.0	6200.0	3200.0	1675.0	225.0
4	11395.0	7150.0	500.0	2945.0	615.0	65.0	140.0

```
[ ]: # confirm dtypes
neighbourhoods.dtypes
```

```
[ ]: neighbourhood    object
      n_id             object
      designation      object
      pop_2016         object
      pop_2011         object
      pop_change       object
      private_dwellings object
      occupied_dwellings object
      pop_dens         object
      area             object
      total_commuters  object
      drive            object
      car_passenger    object
      transit          object
      walk             object
      bike             object
      other            object
      dtype: object
```

5.10 Calculating more columns

Let's fix the population change column and calculate the percentage of commuters who bike.

```
[ ]: neighbourhoods['pop_change'] = neighbourhoods['pop_change'] / 100
      neighbourhoods['pct_bike'] = neighbourhoods['bike'] /
      ↪neighbourhoods['total_commuters']
      neighbourhoods.head()
```

```
[ ]:      neighbourhood n_id    designation pop_2016 pop_2011 \
0      Agincourt North  129  No Designation  29113.0  30279.0
1  Agincourt South-Malvern West  128  No Designation  23757.0  21988.0
2      Alderwood      20  No Designation  12054.0  11904.0
3      Annex          95  No Designation  30526.0  29177.0
4  Banbury-Don Mills    42  No Designation  27695.0  26918.0

      pop_change private_dwellings occupied_dwellings pop_dens area \
0      -0.039      9371.0      9120.0  3929.0  7.41
```

1	0.08	8535.0	8136.0	3034.0	7.83
2	0.013	4732.0	4616.0	2435.0	4.95
3	0.046	18109.0	15934.0	10863.0	2.81
4	0.029	12473.0	12124.0	2775.0	9.98

	total_commuters	drive	car_passenger	transit	walk	bike	other	\
0	11820.0	7155.0	930.0	3350.0	265.0	70.0	45.0	
1	10160.0	6135.0	665.0	2985.0	280.0	35.0	65.0	
2	6045.0	4090.0	355.0	1285.0	195.0	65.0	65.0	
3	14910.0	3290.0	290.0	6200.0	3200.0	1675.0	225.0	
4	11395.0	7150.0	500.0	2945.0	615.0	65.0	140.0	

	pct_bike
0	0.005922
1	0.003445
2	0.010753
3	0.112341
4	0.005704

5.11 merge()ing

The profile are now ready to merge into the bike thefts data!

```
[ ]: thefts_demo = pd.merge(thefts,
                             neighbourhods,
                             how='left',
                             left_on='hood_id',
                             right_on='n_id')
thefts_demo.head()
```

	_id	objectid	event_unique_id	primary_offence	occurrence_date	\
0	1	17744	GO-20179016397	THEFT UNDER	2017-10-03	
1	2	17759	GO-20172033056	THEFT UNDER - BICYCLE	2017-11-08	
2	3	17906	GO-20189030822	THEFT UNDER - BICYCLE	2018-09-14	
3	4	17962	GO-2015804467	THEFT UNDER	2015-05-07	
4	5	17963	GO-20159002781	THEFT UNDER	2015-05-16	

	occurrence_year	occurrence_month	occurrence_dayofweek	\
0	2017	October	Tuesday	
1	2017	November	Wednesday	
2	2018	September	Friday	
3	2015	May	Thursday	
4	2015	May	Saturday	

	occurrence_dayofmonth	occurrence_dayofyear	occurrence_hour	report_date	\
0	3	276	14	2017-10-03	

1	8	312	3	2017-11-08
2	14	257	9	2018-09-17
3	7	127	18	2015-05-14
4	16	136	12	2015-05-16

	report_year	report_month	report_dayofweek	report_dayofmonth	\
0	2017	October	Tuesday	3	
1	2017	November	Wednesday	8	
2	2018	September	Monday	17	
3	2015	May	Thursday	14	
4	2015	May	Saturday	16	

	report_dayofyear	report_hour	division	city	hood_id	\
0	276	18	D22	Toronto	15	
1	312	22	D22	Toronto	15	
2	260	16	D22	Toronto	15	
3	134	14	D22	Toronto	15	
4	136	15	D22	Toronto	15	

	neighbourhoodname	location_type	\
0	Kingsway South (15)	Streets, Roads, Highways (Bicycle Path, Privat...	
1	Kingsway South (15)	Single Home, House (Attach Garage, Cottage, Mo...	
2	Kingsway South (15)	Ttc Subway Station	
3	Kingsway South (15)	Ttc Subway Station	
4	Kingsway South (15)	Ttc Subway Station	

	premises_type	bike_make	bike_model	bike_type	bike_speed	bike_colour	\
0	Outside	GI	ESCAPE 2	OT	7	BLK	
1	House	UNKNOWN	MAKE	NaN	1	BLK	
2	Transit	OT	CROSSTRAIL	MT	24	BLK	
3	Transit	GT	NaN	TO	10	BLKDGR	
4	Transit	GI	NaN	MT	6	RED	

	bike_cost	status	objectid2	\
0	700.0	STOLEN	1	
1	1100.0	RECOVERED	2	
2	904.0	STOLEN	3	
3	400.0	STOLEN	4	
4	600.0	STOLEN	5	

	geometry	neighbourhood	n_id	\
0	{'type': 'Point', 'coordinates': (-79.50655965...	Kingsway South	15	
1	{'type': 'Point', 'coordinates': (-79.50484874...	Kingsway South	15	
2	{'type': 'Point', 'coordinates': (-79.51170915...	Kingsway South	15	
3	{'type': 'Point', 'coordinates': (-79.51170915...	Kingsway South	15	
4	{'type': 'Point', 'coordinates': (-79.51132657...	Kingsway South	15	

	designation	pop_2016	pop_2011	pop_change	private_dwellings	\
0	No Designation	9271.0	9170.0	0.011	3710.0	
1	No Designation	9271.0	9170.0	0.011	3710.0	
2	No Designation	9271.0	9170.0	0.011	3710.0	
3	No Designation	9271.0	9170.0	0.011	3710.0	
4	No Designation	9271.0	9170.0	0.011	3710.0	

	occupied_dwllings	pop_dens	area	total_commuters	drive	car_passenger	\
0	3584.0	3593.0	2.58	3735.0	2210.0	120.0	
1	3584.0	3593.0	2.58	3735.0	2210.0	120.0	
2	3584.0	3593.0	2.58	3735.0	2210.0	120.0	
3	3584.0	3593.0	2.58	3735.0	2210.0	120.0	
4	3584.0	3593.0	2.58	3735.0	2210.0	120.0	

	transit	walk	bike	other	pct_bike
0	1185.0	115.0	30.0	50.0	0.008032
1	1185.0	115.0	30.0	50.0	0.008032
2	1185.0	115.0	30.0	50.0	0.008032
3	1185.0	115.0	30.0	50.0	0.008032
4	1185.0	115.0	30.0	50.0	0.008032


5.12 Grouping and plotting

With the datasets joined, we can aggregate and plot the data. We can start using statistical methods, like `corr()` to check for relationships between variables as well.

```
[ ]: thefts_2016_grouped = (thefts_demo
                             .query('occurrence_year == 2016')
                             .groupby(['neighbourhood']))
```


```
[ ]: # neighbourhood populations are skewed
neighbourhoods.query('neighbourhood != "City of Toronto")['pop_2016'].
    ↪plot(kind='hist')
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe3f7e297d0>
```




```
[ ]: # so are the % of commuters who bike to work
neighbourhoods.query('neighbourhood != "City of Toronto"')['pct_bike'].
    plot(kind='hist')
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe3f7b61e50>
```




```
[ ]: # as are thefts
thefts_2016_grouped.size().plot(kind='hist')
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe3f7b5bc90>
```




```
[ ]: # thefts counts vs population
(thefts_2016_grouped
 .agg(total_thefts=('_id', 'count'),
      pop=('pop_2016', 'median'),
      pct_bike=('pct_bike', 'mean'))
 .reset_index()
 .plot(kind='scatter', y='total_thefts', x='pop'))
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe3f7a3bcd0>
```



```
[ ]: (thefts_2016_grouped
      .size()
      .sort_values(ascending=False)
      .head(10)
      .plot(kind='bar', title='Reported Bike Thefts, 2016'))
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe3f79c4b90>
```



```
[ ]: # a quick correlation check
(thefts_2016_grouped
 .agg(total_thefts=('_id', 'count'),
      pop=('pop_2016', 'median'),
      dens=('pop_dens', 'median'),
      pct_bike=('pct_bike', 'mean'))
 .corr('spearman'))
```

```
[ ]:
total_thefts    total_thefts      pop      dens  pct_bike
total_thefts      1.000000  0.267761  0.485556  0.651319
```

pop	0.267761	1.000000	0.020082	-0.222565
dens	0.485556	0.020082	1.000000	0.605242
pct_bike	0.651319	-0.222565	0.605242	1.000000

```
[ ]: thefts_demo.to_csv('/content/data/bike_thefts_joined.csv', index=False)
      neighbourhoods.to_csv('/content/data/neighbourhoods.csv', index=False)
```

6 References

6.0.1 Programming

- pandas development team. *API reference*. <https://pandas.pydata.org/pandas-docs/stable/reference/index.html>
- pandas development team. *User guide*. https://pandas.pydata.org/pandas-docs/stable/user_guide/index.html
- *Python strptime cheatsheet*. <https://strftime.org/>

6.0.2 Data Sources

- Open Data Toronto. *Neighbourhood Profiles*. <https://open.toronto.ca/dataset/neighbourhood-profiles/>
- Open Data Toronto. *TTC Subway Delay Data*. <https://open.toronto.ca/dataset/ttc-subway-delay-data/>
- Open Data Toronto. *Bicycle Thefts*. <https://open.toronto.ca/dataset/bicycle-thefts/>